# Probit Regression Function

## Applications in Government & Politics

Elliot Hipski, Rohan, Brandon, Arnav

April 2026

## Table of Contents

1. Prerequisite Knowledge
2. General Explanation
3. Mathematical Explanation
4. Implementation
5. Comparison to Other Models
6. Citations and Appendix

# Topic 1 | Prerequisite Knowledge

## Bernoulli Distribution Variance

Used for modeling binary outcomes. The variance is defined as $\text{Var}(X) = E[X^2] - (E[X])^2 = p - p^2 = p(1-p)$

## General Linearized Model

A *General Linearized Model* is a model that applies a linear function to a dataset, with the general rules of mapping the input to the output y, and mapping it according to a **Link Function**.

## Gradient

A vector of partial derivatives. It is used in mapping the log likelihood function of the linear model. For this case, we use the weights as slopes and apply the transformation $X^T \cdot W$, where $W$ is the weights matrix. This is used in solving for the **Hessian**

## Hessian

A matrix of 2nd order partial derivatives. For this case, however, we take the negative Expected Value of the hessian, which simplifies the matrix math. Using the previously defined weights matrix, a formula for solving the Expected Hessian becomes (2):

# Topic 2 | General Explanation

## The Probit Mechanism

- **Distribution Assumption**: Assumes an underlying normal distribution for the latent variable.
- **Inputs**: Utilizes the CDF ($\Phi$) and PDF ($\phi$) of the standard normal distribution.
- **Predictor Matrix**: Uses a linear predictor $X\beta$, where $X$ contains a column of 1s (intercept) and a column of observations.
- **Score Calculation**: Calculated using the difference between $y$ variables and the normal PDF, scaled by the ratio of the PDF to the Bernoulli variance.
- **Iteration**: The system solves for $\beta$ by using the Hessian and Gradient to update variables until convergence (difference $<$ tolerance).

# Topic 3 | Mathematical Explanation

# Topic 4 | Implementation

```r
myprobit = function(X, y, max_iter, tolerance, verbose=FALSE){

  m = ncol(X)
  beta <- rep(0, m)

  for (i in 1:max_iter) {

    eta <- X %*% beta

    phat <- pnorm(eta)
    qhat <- dnorm(eta)

    ratio <- qhat / (phat * (1 - phat))
    score <- (y - phat) * ratio
    g <- t(X) %*% score

    w <- (qhat^2) / (phat * (1 - phat))
    H <- -t(X) %*% (X * as.vector(w))

    copy_beta = beta
    beta <- beta - solve(H, g)

    if (verbose){
      cat("Iteration", i, ":\n")
      j=1
      for (b in beta){
        cat("Beta", j, " = ", beta[j], "\n")
        j=j+1
      }
    }
    if ((sum(beta - copy_beta)^2) < tolerance) {
        return(beta)
        break
      }
  }
}
```

```r
X <- model.matrix(am ~ mpg, data = mtcars)
y <- mtcars$am
myprobit(X,y,32, 0.00009, FALSE)
```

```
##                   [,1]
## (Intercept) -3.9100663
## mpg          0.1822605
```

# Solving the System

## Optimization Methods

- There are 2 main methods of creating a model for probit regression: ### 1. Newton-Raphson method We used a method similar to Newton Raphson, but rather than use the Hessian, the Expected Hessian was used, which is used for IRLS and is simpler to calculate. ### 2. Iteratively Reweighted Least Squares / Fisher Scoring
- These methods can be shown to be mathematically the same(7) ## Testing

# Topic 5 | Comparison to Other Models

## Motivation

Several methods exist for modeling binary outcomes. This section compares probit regression against three alternatives:

1. **Linear Regression (OLS)** — the standard baseline for continuous outcomes
2. **Logistic Regression** — the most widely used GLM for binary data
3. **Lasso Regression** — a penalized method that performs automatic variable selection

The goal is to understand the practical and theoretical differences between these approaches, and to identify when probit regression is the appropriate choice.

## Data: Donner Party Survival

We use the **Donner Party survival dataset** from the `Sleuth3` package. The binary outcome is survival status (1 = Survived, 0 = Died), with Age and Sex as predictors.

```r
library(Sleuth3); library(dplyr); library(ggplot2)
library(glmnet);  library(caret)

set.seed(42)

# recode response and sex
donner <- case2001 %>%
  mutate(
    mystatus = if_else(Status == "Survived", 1, 0),
    male     = if_else(Sex == "Male", 1, 0)
  )

head(donner[, c("Age", "Sex", "Status", "mystatus")], 4)
```

## Data Preview

```
  Age    Sex   Status mystatus
1  23   Male     Died        0
2  40 Female Survived        1
3  40   Male Survived        1
4  30   Male     Died        0
```

## Linear Regression on a Binary Response

OLS fits:

$$\hat{E}(Y \mid \mathbf{X}) = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \cdots + \hat{\beta}_p X_p$$

When $Y \in \{0, 1\}$, this is the **Linear Probability Model (LPM)**. Since $\mathbf{X} \in \mathbb{R}$, fitted values are not constrained to $[0, 1]$, which means the model can produce predicted probabilities below 0 or above 1. This violates the basic requirement for a valid probability model and is the primary motivation for using a link function in GLMs.

# OLS on a Binary Response — Code

```r
# fit OLS treating the binary outcome as continuous
lpm_fit <- lm(mystatus ~ Age + male, data = donner)

# store fitted values as predicted probabilities (not bounded to [0,1])
donner <- donner %>%
  mutate(pi_lpm = predict(lpm_fit))
```

# OLS on a Binary Response — Plot Code

```
ggplot(donner, aes(x = Age, y = mystatus)) +
  geom_point(aes(color = Status), size = 2, alpha = 0.7) +
  geom_line(aes(y = pi_lpm, group = male,
                linetype = factor(male, labels = c("Female","Male"))),
            color = "darkred", linewidth = 1) +
  geom_hline(yintercept = c(0, 1), linetype = "dotted", color = "gray40") +
  labs(title = "Linear Probability Model: Predictions Can Leave [0, 1]",
       y = "P(Survived)", x = "Age", linetype = "Sex",
       caption = "Dotted lines mark the valid [0,1] probability bounds") +
  theme_bw(base_size = 12)
```

# OLS on a Binary Response — Plot

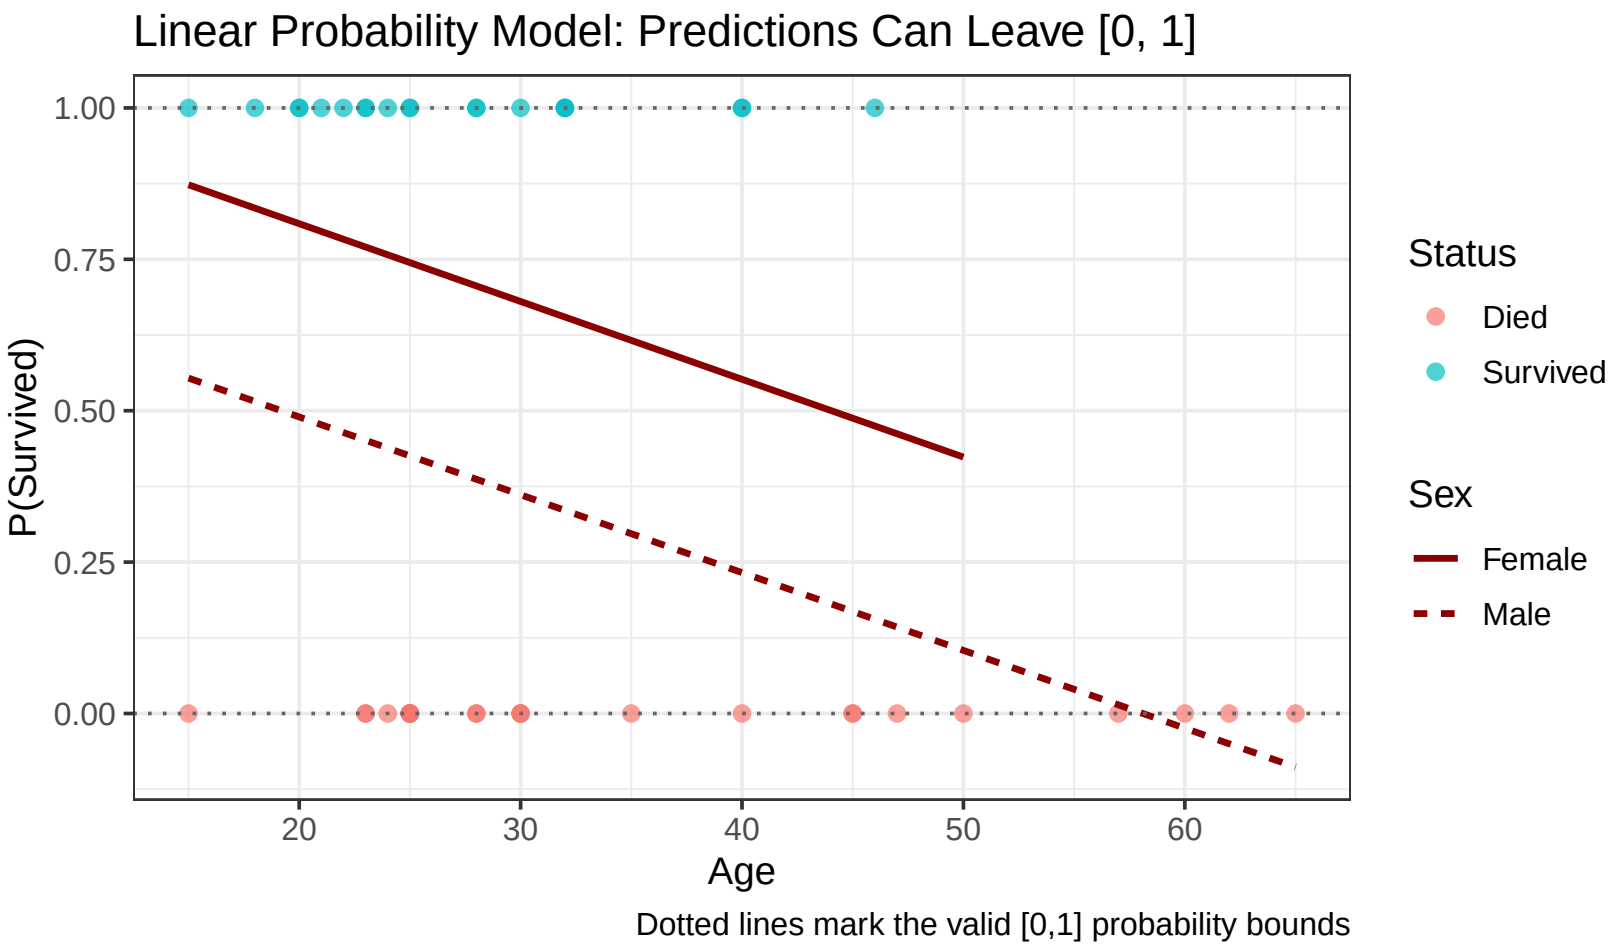

# OLS Predictions Outside [0, 1]

```r
lpm_preds <- predict(lpm_fit)

cat("Range of OLS predictions:", round(range(lpm_preds), 4), "\n")
```

```
Range of OLS predictions: -0.0885 0.8731
```

```r
cat("Predictions below 0:     ", sum(lpm_preds < 0), "\n")
```

```
Predictions below 0:      3
```

```r
cat("Predictions above 1:     ", sum(lpm_preds > 1), "\n")
```

```
Predictions above 1:      0
```

```r
# misclassification at 0.5 threshold
lpm_class <- ifelse(lpm_preds > 0.5, 1, 0)
cat("OLS misclassification rate:",
    round(mean(lpm_class != donner$mystatus), 4), "\n")
```

```
OLS misclassification rate: 0.2444
```

Even on this small dataset, OLS produces invalid probabilities. Linear regression is not appropriate for binary outcomes.

## Logistic Regression

Logistic regression uses the **logit link**:

$$\eta = \text{logit}(\pi) = \ln\left[\frac{\pi}{1-\pi}\right] = \mathbf{x}^T\boldsymbol{\beta} \quad \Rightarrow \quad \hat{\pi} = \frac{e^\eta}{1+e^\eta}$$

This assumes errors follow a logistic distribution.
Probit regression uses the **inverse standard normal CDF**:

$$\eta = \Phi^{-1}(\pi) = \mathbf{x}^T\boldsymbol{\beta} \quad \Rightarrow \quad \hat{\pi} = \Phi(\mathbf{x}^T\hat{\boldsymbol{\beta}})$$

This assumes errors follow a standard normal distribution. Both are valid GLMs for binary outcomes — they differ only in their distributional assumption.

## Link Functions — Code

```r
eta_seq <- seq(-4, 4, length.out = 300)

df_links <- data.frame(
  eta    = rep(eta_seq, 2),
  pi     = c(plogis(eta_seq),  # inverse logit
             pnorm(eta_seq)),  # inverse probit (std normal CDF)
  Method = rep(c("Logistic (logit link)", "Probit (normal CDF)"), each = 300)
)
```

# Link Functions — Plot Code

```r
ggplot(df_links, aes(x = eta, y = pi, color = Method, linetype = Method)) +
  geom_line(linewidth = 1.3) +
  labs(title = "Inverse Link Functions: Logistic vs. Probit",
       x = expression(eta ~ "(Linear Predictor)"),
       y = expression(hat(pi) ~ "(Predicted Probability)")) +
  scale_color_manual(values = c("steelblue", "firebrick")) +
  theme_bw(base_size = 12) + theme(legend.position = "bottom")
```

# Link Functions — Plot

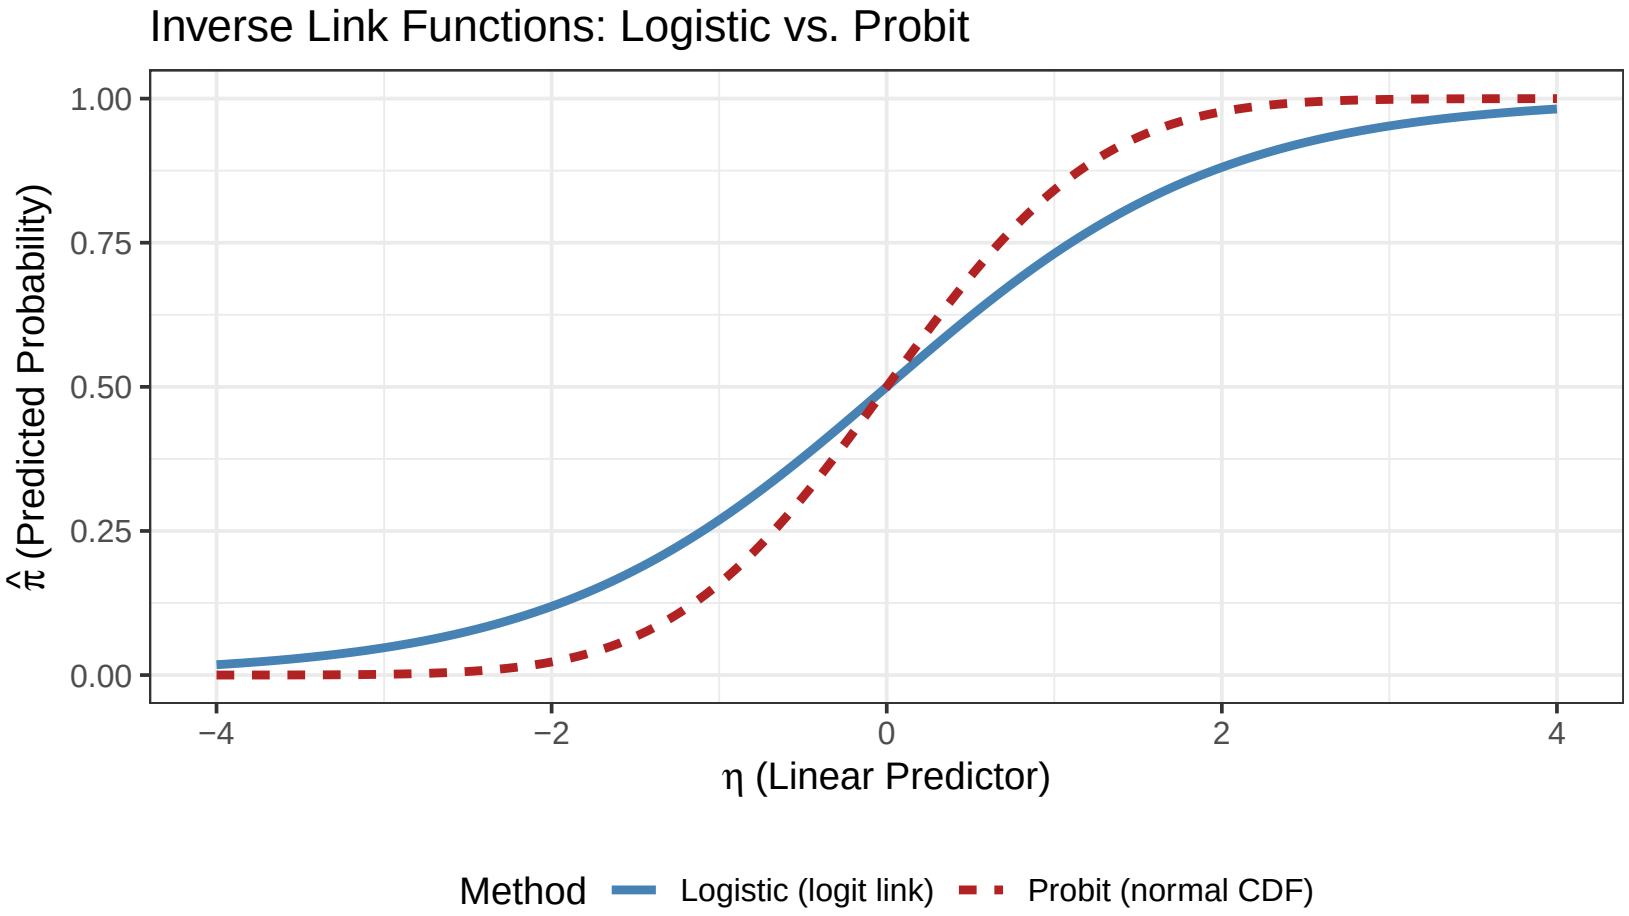


The probit curve has lighter tails — it approaches 0 and 1 faster than the logistic curve.

# Fitting Both Models — Code

```r
# fit logistic and probit
logit_fit  <- glm(mystatus ~ Age + male, data = donner,
                  family = binomial(link = "logit"))

probit_fit <- glm(mystatus ~ Age + male, data = donner,
                  family = binomial(link = "probit"))

# compare coefficients
round(cbind(
  "OLS Coef"    = coef(lpm_fit),
  "Logit Coef"  = coef(logit_fit),
  "Probit Coef" = coef(probit_fit)
), 4)
```

## Fitting Both Models — Output

```
            OLS Coef Logit Coef Probit Coef
(Intercept)   1.0658     3.2304      1.9173
Age          -0.0128    -0.0782     -0.0457
male         -0.3192    -1.5973     -0.9583
```

The probit coefficients are smaller by a factor of roughly 1.6 to 1.8. The logistic distribution has variance $\pi^2/3 \approx 3.29$, while the standard normal has variance 1. The scales differ, but the substantive conclusions are the same.

## Predicted Probabilities — Code

```r
donner <- donner %>%
  mutate(
    pi_logit  = predict(logit_fit,  type = "response"),
    pi_probit = predict(probit_fit, type = "response")
  )
```

# Predicted Probabilities — Plot Code

```r
ggplot(donner, aes(x = pi_logit, y = pi_probit, color = Status)) +
  geom_point(size = 2.5, alpha = 0.8) +
  geom_abline(slope = 1, intercept = 0, linetype = "dashed", color = "gray40") +
  labs(title = "Logistic vs. Probit: Predicted Probabilities (Donner Data)",
       x = "Logistic Predicted Probability",
       y = "Probit Predicted Probability") +
  theme_bw(base_size = 12)
```

# Predicted Probabilities — Plot

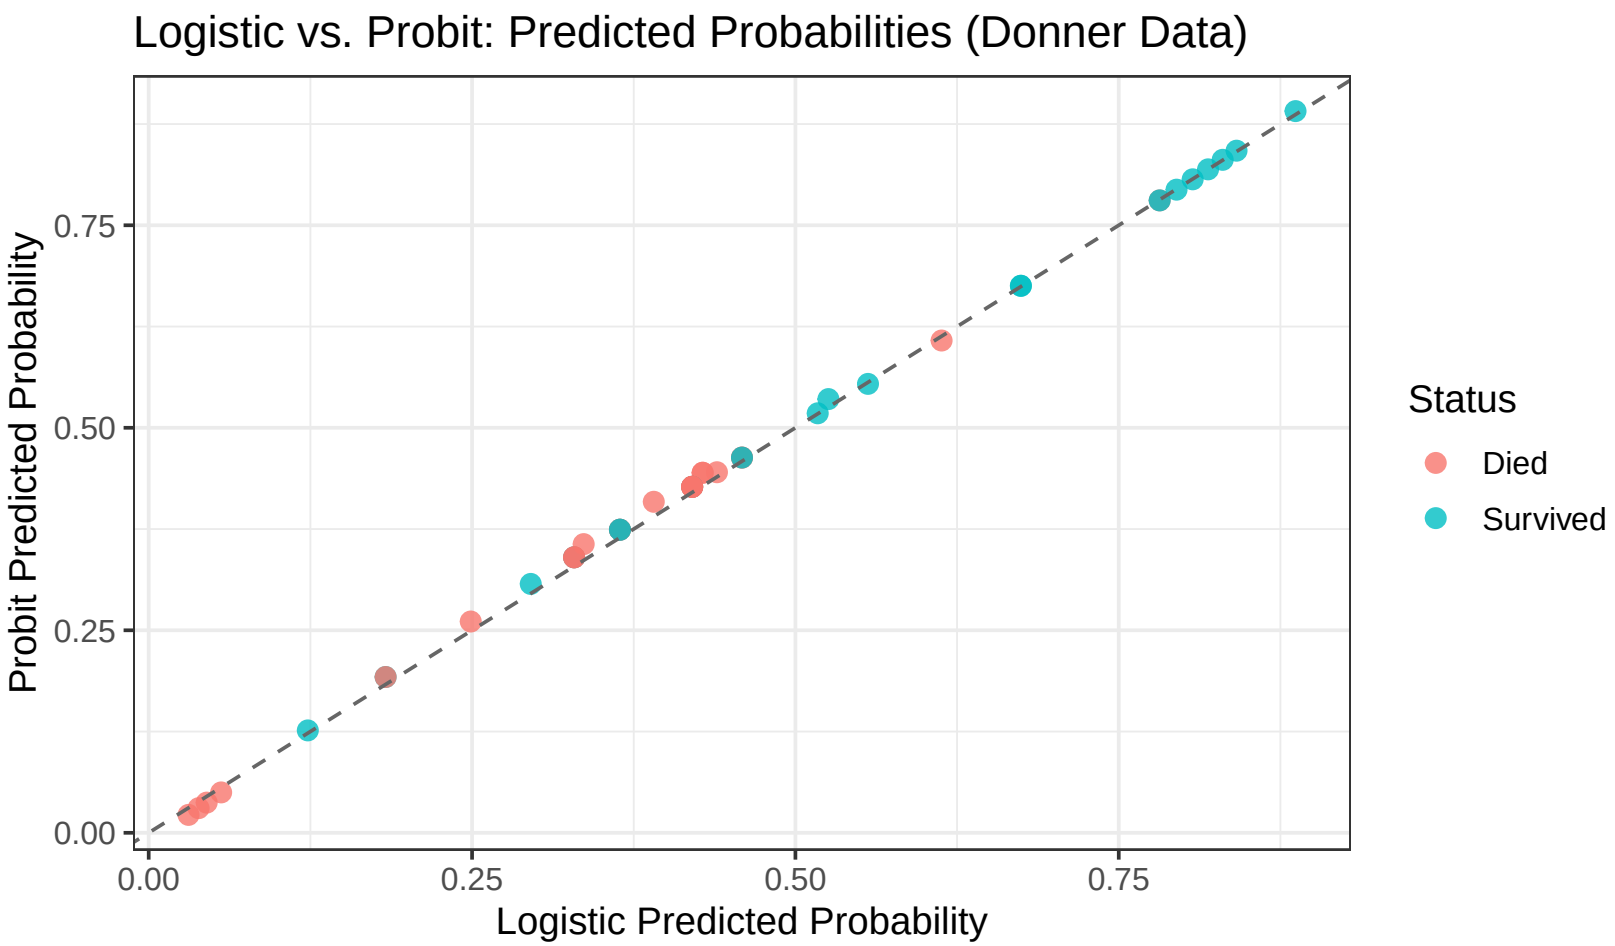


The predicted probabilities from both models are nearly identical, confirming that the choice of link function has minimal practical impact on fitted values.

# Model Fit Metrics — Code

```r
cat(sprintf("%-28s %8s %8s\n", "Metric", "Logistic", "Probit"))
cat(sprintf("%-28s %8.3f %8.3f\n", "AIC",
            AIC(logit_fit), AIC(probit_fit)))
cat(sprintf("%-28s %8.3f %8.3f\n", "Residual Deviance",
            deviance(logit_fit), deviance(probit_fit)))
cat(sprintf("%-28s %8.3f %8.3f\n", "Null Deviance",
            logit_fit$null.deviance, probit_fit$null.deviance))

log_err  <- mean((donner$pi_logit  > 0.5) != donner$mystatus)
prob_err <- mean((donner$pi_probit > 0.5) != donner$mystatus)
cat(sprintf("%-28s %8.4f %8.4f\n", "Misclassification Rate",
            log_err, prob_err))
```

## Model Fit Metrics — Output

| Metric | Logistic | Probit |
|---|---|---|
| AIC | 57.256 | 57.283 |
| Residual Deviance | 51.256 | 51.283 |
| Null Deviance | 61.827 | 61.827 |
| Misclassification Rate | 0.2222 | 0.2222 |

The two models produce nearly identical fit across every metric. AIC differs by less than 0.03, and misclassification rates are equal. The choice between logistic and probit is rarely driven by model performance — it is almost always a question of convention, interpretability, or the underlying distributional assumption.

## When to Prefer Probit

**1. Latent variable framework.** Assume an unobserved variable $Y^* = \mathbf{x}^T\beta + \varepsilon$ where $\varepsilon \sim N(0,1)$. We observe $Y = 1$ if $Y^* > 0$. This is the probit model exactly — a standard framework in economics and political science.

**2. Tail behavior.** The normal CDF approaches 0 and 1 more rapidly than the logistic distribution. Probit is preferable when probabilities near the extremes are of scientific interest.

**3. Field convention.** Probit is the standard choice in economics, political science, and related disciplines. For government and political science applications, probit is the conventional model.

**4. Marginal effects.** Probit coefficients do not have an odds-ratio interpretation, but average marginal effects provide a clear and directly interpretable quantity.

# Marginal Effects

Probit coefficients have no direct probability interpretation on their own. The standard approach is to compute **Average Marginal Effects (AME)**:

$$\text{AME}_j = \frac{1}{n} \sum_{i=1}^{n} \phi(\mathbf{x}_i^T \hat{\boldsymbol{\beta}}) \cdot \hat{\beta}_j$$

where $\phi(\cdot)$ is the standard normal PDF.

```r
library(margins)

# compute AMEs
ame <- margins(probit_fit)
summary(ame)[, c("factor", "AME", "SE", "z", "p")]
```

```
 factor     AME     SE       z      p
    Age -0.0148 0.0057 -2.6113 0.0090
   male -0.3109 0.1188 -2.6160 0.0089
```

This gives the change in predicted probability for a one-unit change in each predictor, averaged across all observations.

## Lasso Regression

The logistic or probit objective function can be **penalized** to perform variable selection:

$$\hat{\boldsymbol{\beta}} = \text{argmin}_{\boldsymbol{\beta}} \left\{ -\frac{2}{n}\ell(\boldsymbol{\beta}) + \lambda \sum_{j=1}^{p} |\beta_j| \right\}$$

The $\ell_1$ penalty shrinks small coefficients toward zero and eliminates irrelevant predictors entirely when $\lambda$ is large enough. This is particularly useful in high-dimensional settings such as political science data with many demographic and behavioral predictors.

## Lasso: Fitting the Model — Code

```r
# glmnet requires a plain matrix, not a model formula
X_mat <- model.matrix(mystatus ~ Age + male, data = donner)[, -1]
y_vec <- donner$mystatus

# fit lasso path over 100 values of lambda (alpha=1 selects lasso over ridge)
lasso_path <- glmnet(X_mat, y_vec, family = "binomial", alpha = 1)
```

# Lasso: Coefficient Path

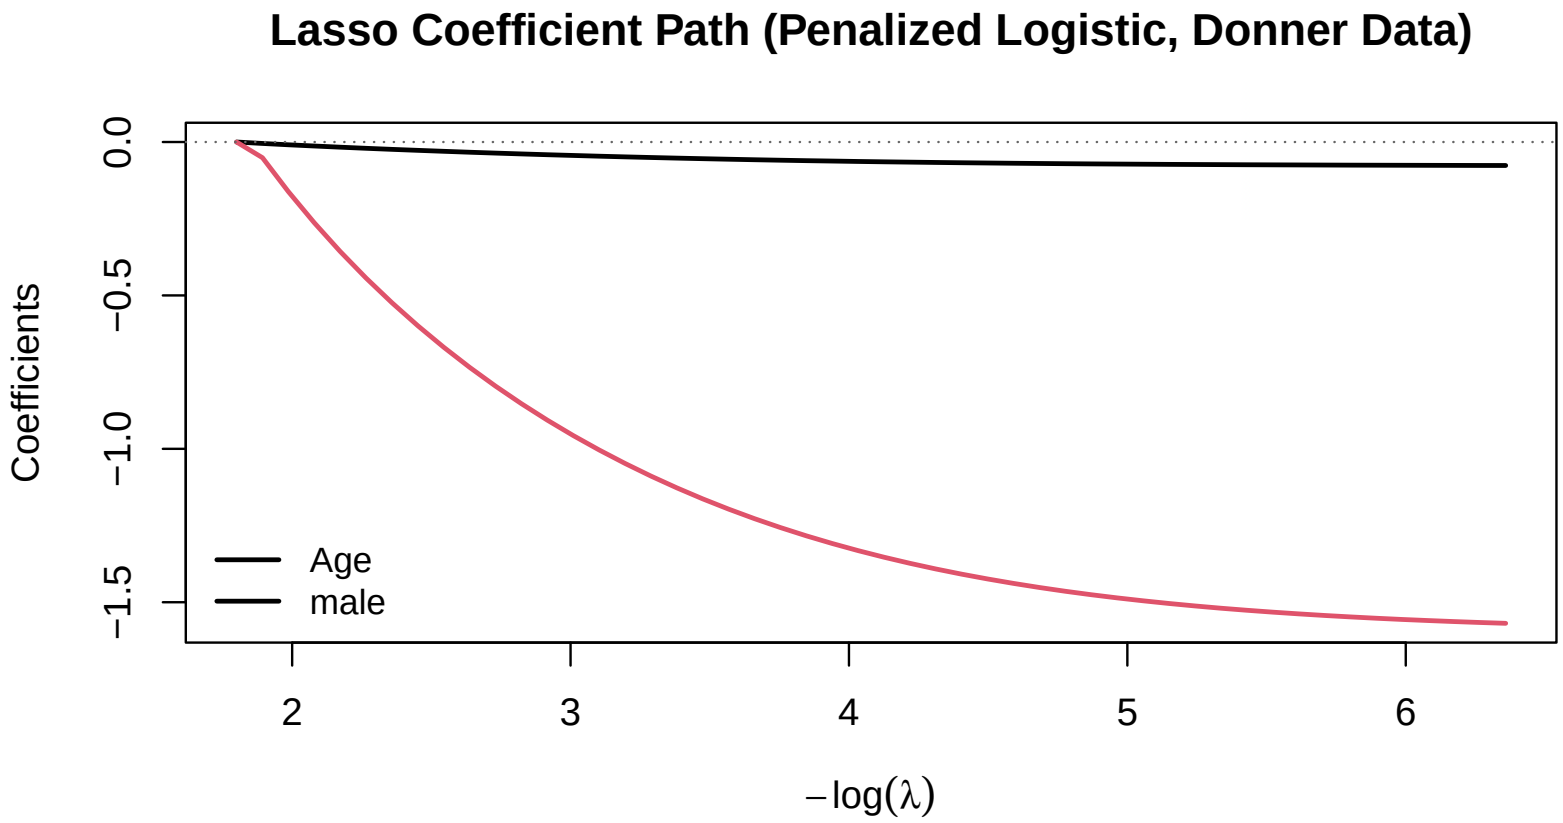


As $\lambda$ increases (moving left on the x-axis), both coefficients shrink toward zero. The `male` coefficient shrinks more slowly, indicating it is the stronger predictor of survival. `Age` enters the model quickly but with a smaller effect.

# Lambda Selection via Cross-Validation — Code

```r
set.seed(42)
# 10-fold CV to select lambda; type.measure="class" minimizes misclassification
cv_lasso <- cv.glmnet(X_mat, y_vec,
                      family       = "binomial",
                      alpha        = 1,
                      nfolds       = 10,
                      type.measure = "class")

# lambda.min: lowest CV error; lambda.1se: most regularized model within 1 SE
cat("Lambda (min CV error): ", round(cv_lasso$lambda.min, 5), "\n")
```

```
Lambda (min CV error):  0.02822
```

```r
cat("Lambda (1 SE rule):    ", round(cv_lasso$lambda.1se, 5), "\n")
```

```
Lambda (1 SE rule):     0.06518
```

## Lambda Selection via Cross-Validation — Plot

10-fold CV selects $\lambda_{\text{min}}$ as the value minimizing misclassification error, and $\lambda_{\text{1SE}}$ as the largest $\lambda$ within one standard error of that minimum — the more parsimonious choice. Error stabilizes well before the full-model limit, indicating that a small penalty is sufficient here.

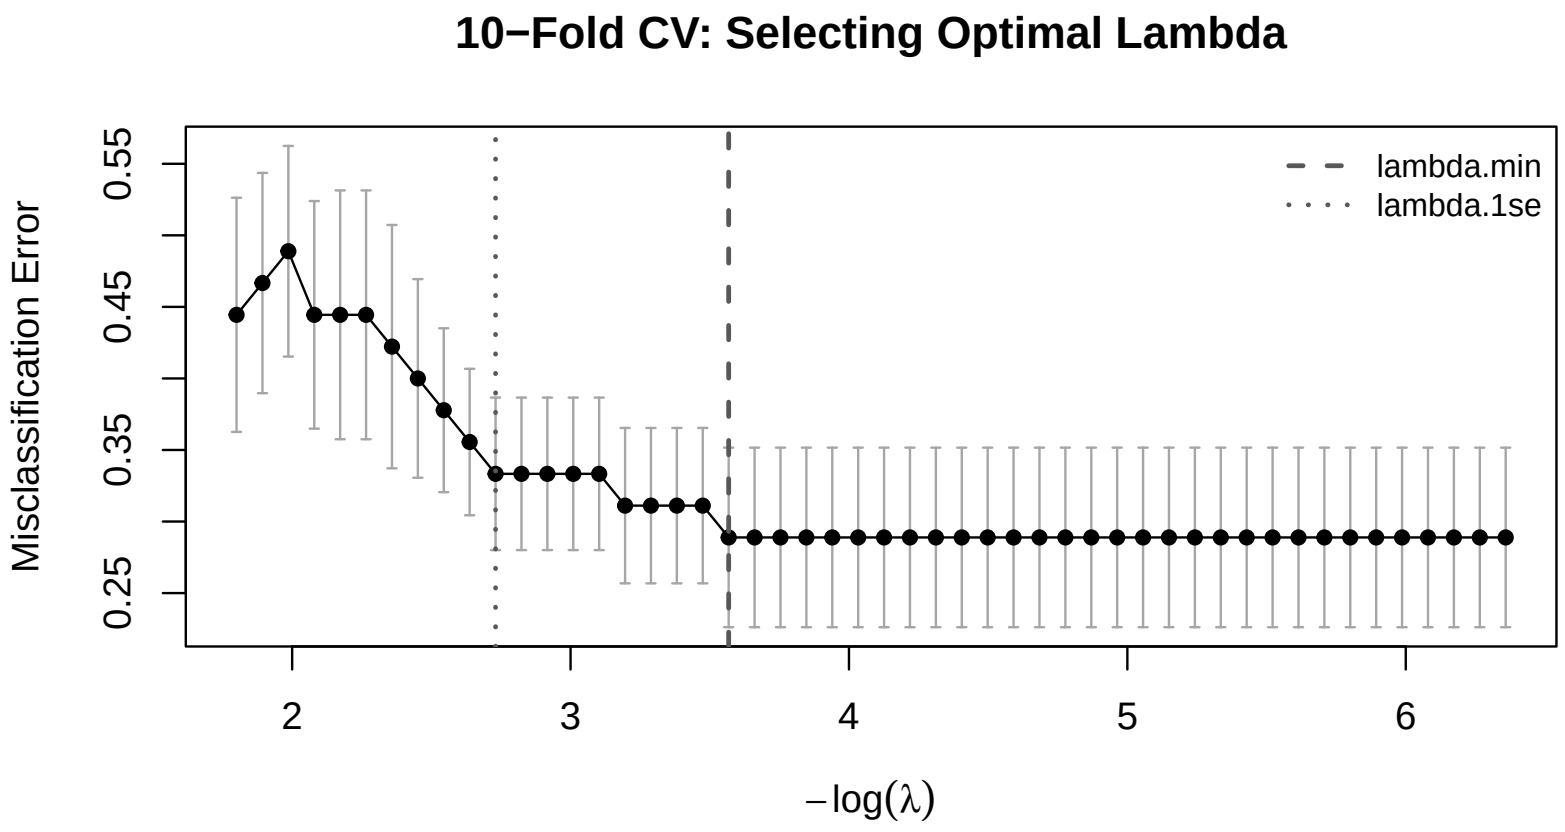

# Lasso Final Coefficients

```r
cat("Lasso coefficients at lambda.min:\n")
```

```
Lasso coefficients at lambda.min:
```

```r
round(as.matrix(coef(cv_lasso, s = "lambda.min")), 5)
```

```
            lambda.min
(Intercept)    2.30441
Age           -0.05615
male          -1.19653
```

```r
cat("\nLasso coefficients at lambda.1se:\n")
```

```

Lasso coefficients at lambda.1se:
```

```r
round(as.matrix(coef(cv_lasso, s = "lambda.1se")), 5)
```

```
            lambda.1se
(Intercept)    1.43295
Age           -0.03615
male          -0.79566
```

A dot (.) indicates a coefficient shrunk exactly to zero — that predictor has been

# All Four Methods — Code

```
ggplot(donner, aes(x = Age)) +
  geom_point(aes(y = mystatus, color = Status), alpha = 0.5, size = 1.5) +
  geom_line(aes(y = pi_lpm, group = male), color = "gray30",
            linewidth = 0.9, linetype = "dotdash") +
  geom_smooth(aes(y = mystatus), method = "glm",
              method.args = list(family = binomial(link = "logit")),
              se = FALSE, color = "steelblue", linewidth = 1) +
  geom_smooth(aes(y = mystatus), method = "glm",
              method.args = list(family = binomial(link = "probit")),
              se = FALSE, color = "firebrick", linewidth = 1, linetype = "dashed") +
  labs(title = "Predicted Probabilities: All Four Methods (Donner Data)",
       y = "P(Survived)", x = "Age",
       caption = "Blue = Logistic  |  Red = Probit  |  Gray = OLS  |  Lasso approx. Logistic") +
  theme_bw(base_size = 12)
```

# All Four Methods — Plot

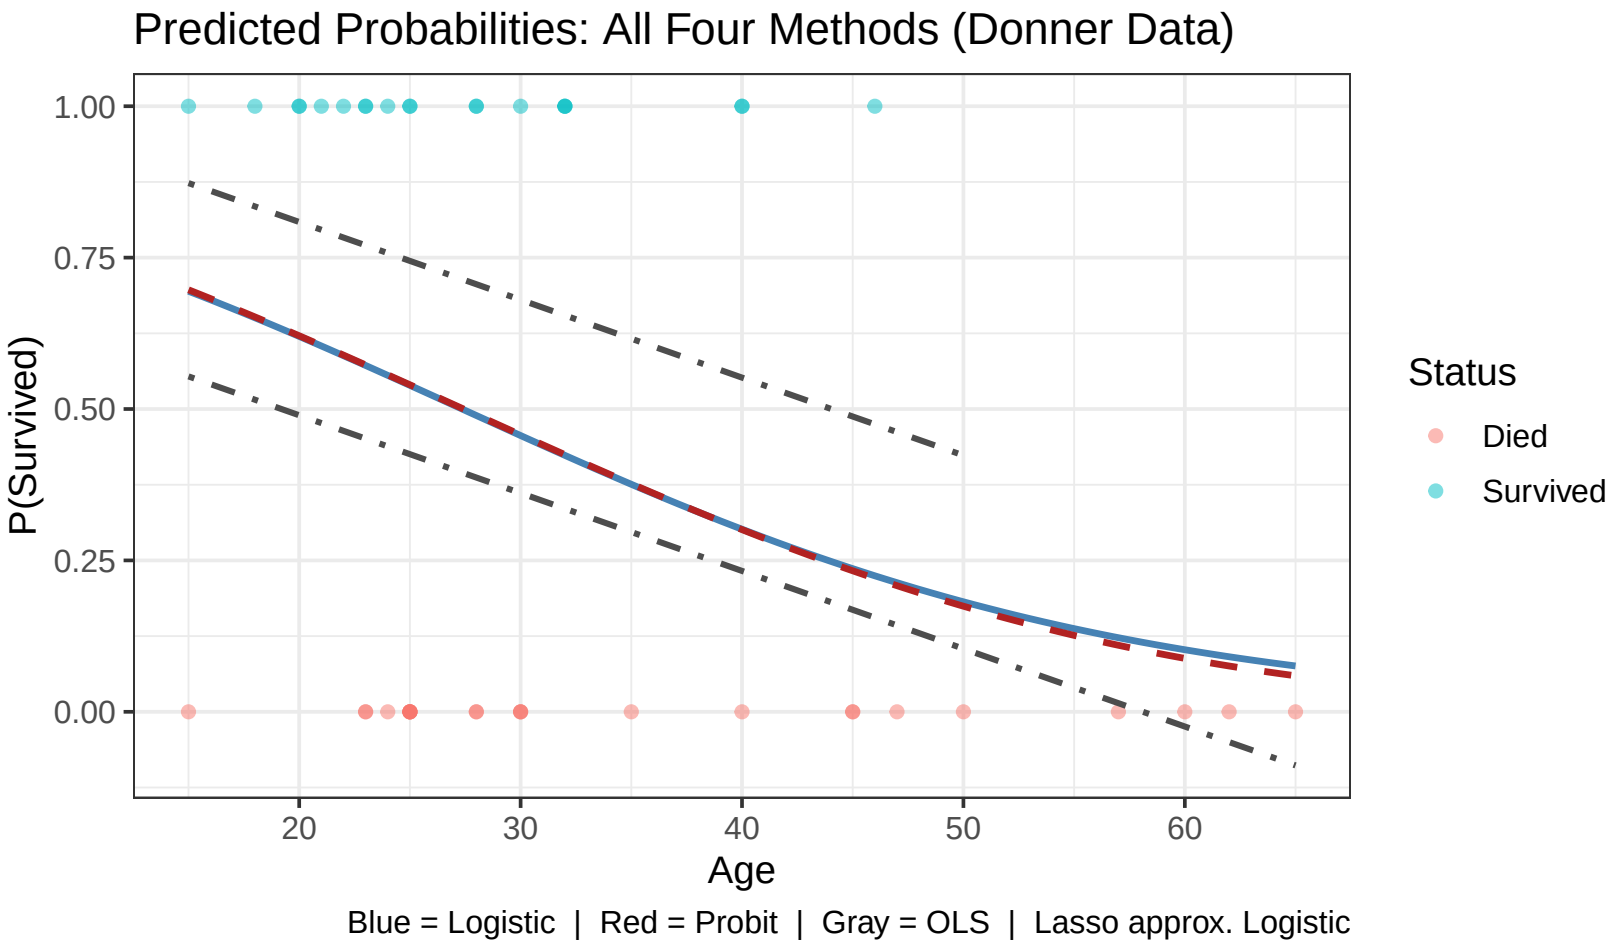

# Confusion Matrix Comparison — Code

```r
# classify each observation using a 0.5 probability threshold
threshold <- 0.5

results <- donner %>%
  mutate(
    # convert predicted probabilities to binary class labels
    pred_lpm    = factor(ifelse(pi_lpm    > threshold, 1, 0)),
    pred_logit  = factor(ifelse(pi_logit  > threshold, 1, 0)),
    pred_probit = factor(ifelse(pi_probit > threshold, 1, 0)),
    # lasso: get response-scale predictions at the CV-selected lambda
    pred_lasso  = factor(ifelse(
      predict(cv_lasso, newx = X_mat, s = "lambda.min",
              type = "response")[, 1] > threshold, 1, 0)),
    actual = factor(mystatus)
  )
```

# Confusion Matrix Comparison — Accuracy Code

```r
# helper: proportion of correctly classified observations
acc <- function(pred, actual) round(mean(pred == actual), 4)

cat(sprintf("%-20s %s\n", "Method",    "Accuracy"))
cat(sprintf("%-20s %s\n", "OLS (LPM)", acc(results$pred_lpm,    results$actual)))
cat(sprintf("%-20s %s\n", "Logistic",  acc(results$pred_logit,  results$actual)))
cat(sprintf("%-20s %s\n", "Probit",    acc(results$pred_probit, results$actual)))
cat(sprintf("%-20s %s\n", "Lasso",     acc(results$pred_lasso,  results$actual)))
```

## Confusion Matrix Comparison — Output

```
Method                  Accuracy

OLS (LPM)               0.7556

Logistic                0.7778

Probit                  0.7778

Lasso                   0.7556
```

Logistic and probit tie at 77.78%, while OLS and lasso both fall to 75.56%. The bounded probability models classify slightly better here, though on a dataset of 45 observations the difference is small. Probit and logistic are effectively equivalent in predictive accuracy on this data.

## Probit Confusion Matrix

```
          Reference
Prediction  0  1
         0 23  8
         1  2 12

Accuracy             : 0.7778
Sensitivity          : 0.6000
Specificity          : 0.9200
Pos Pred Value       : 0.8571
Neg Pred Value       : 0.7419
Kappa                : NA
```

The probit model is considerably better at identifying deaths (specificity = 0.92) than survivors (sensitivity = 0.60). In the Donner Party data, older males had substantially lower survival odds, and the model captures that pattern well at the majority class. The imbalance between sensitivity and specificity is typical when the outcome is unequally distributed.

## Model Comparison Summary

| Criterion | OLS (LPM) | Logistic | Probit | Lasso |
|---|---|---|---|---|
| Predictions in [0, 1]? | No | Yes | Yes | Yes |
| Link function | Identity | Logit | $\Phi^{-1}$ (normal) | Logit or Probit |
| Error distribution | Normal | Logistic | Normal | Logistic |
| Estimation method | OLS (closed form) | MLE (IRLS) | MLE (Fisher Scoring) | MLE + $\ell_1$ penalty |

## Model Comparison Summary

| Criterion | OLS (LPM) | Logistic | Probit | Lasso |
|---|---|---|---|---|
| Coefficient interpretation | $\Delta\hat{y}$ | Log-odds / Odds ratio | Marginal effects (AME) | Log-odds (shrunk) |
| Variable selection? | No | No | No | Yes |
| Tuning required? | No | No | No | Yes ($\lambda$) |
| Preferred when... | Never | General / odds needed | Econ, politics, latent variable | Many predictors |
| Inference (p-values, CIs)? | Yes | Yes | Yes | Not standard |

## Key Takeaways

- **Linear regression (OLS)** is not appropriate for binary outcomes. Predicted values are not bounded to $[0, 1]$ and the model is misspecified.
- **Logistic and probit regression** produce nearly identical predicted probabilities and model fit. Coefficients differ in scale by approximately 1.6x due to differing error distributions, but substantive conclusions are the same.
- **Probit is preferred** in economics and political science due to the latent variable framework and field convention. Coefficients are interpreted via average marginal effects.
- **Logistic regression is preferred** when odds ratios are the natural effect measure.
- **Lasso** extends either model with an $\ell_1$ penalty for automatic variable selection. Standard p-values and confidence intervals are not available after lasso selection.
- For government and political science applications, probit regression is the standard and appropriate choice.

# References

## Citations

**Comparison Section (Rohan)**


- Lucero, C. Lecture 16: Classification & Generalized Linear Models. CMDA-4654, Virginia Tech.
- Urban Institute. Logit and Probit Regression. https://www.urban.org/research/data-methods/data-analysis/quantitative-data-analysis/impact-analysis/regression-techniques/logit-and-probit-regression
- MetricGate. Probit vs. Logit Models for Binary Outcomes. https://metricgate.com/blogs/probit-vs-logit-regression/
- Abbott, M.G. ECON 452 Note 15: Marginal Effects in Probit Models. Queen's University. https://econ.queensu.ca/faculty/abbott/econ452/452note15.pdf
- Friedman, J., Hastie, T., Tibshirani, R. (2010). Regularization Paths for Generalized Linear Models via Coordinate Descent. *Journal of Statistical Software*, 33(1). https://glmnet.stanford.edu
- R Packages: `glmnet`, `caret`, `margins`, `Sleuth3`, `dplyr`,